

Transformer Riffage

Vishakh Begari

Independent Researcher

Gauting, 82131, Germany

DJENTLEVIBE32168@GMAIL.COM

Editor: Adam Smasher

Abstract

Ever since the electrification of guitar, there has been an explosion in rhythmic complexity, tone experimentation and playing techniques. Music generation using Transformers ? has generally been used for creating mostly musical notes. In this paper, a new approach is proposed to capture the diverse expressions commonly featured in modern metal music. A new tokenisation method encapsulates the various guitar techniques and a sequence level attention component which tries to capture the segmentation between different musical phrases. An inference method is implemented that provides meaningful validation without human evaluation.

Keywords: Transformers, Music, Generative

1 Introduction

A *riff* is a short, repeated motif or figure in the melody or accompaniment of a musical composition. ? This type of composition is particularly prevalent in rock, metal, blues, funk and jazz. It is often the main feature that gives character and structure to a song. Recent developments in Transformers architecture have inspired many to generate sequences by treating musical notes as symbolic language.

Musical elements like notes, timestamps, and velocities are converted into discrete tokens similar to text processing. ? A widely adopted method for representing musical notes digitally is the MIDI format, which encodes pitch, duration, and dynamics. Using such a format for musical expression has limitations like resolution and bandwidth constraints ? when compared to richer representations like audio recordings and detailed scores. Modern guitar riffs frequently employ intense and abrupt expressive techniques, which contribute significantly to its overall identity. To accurately represent the nuanced expressive elements of modern music, a score-based format is preferable over simpler encodings like MIDI. A comprehensive list of these expressive elements is provided in Appendix ???. Generating modern music demands capturing not only note patterns but also the associated expressions.

2 Encoding

The original Transformer model ? used a vocabulary of tokens derived from the input language corpus, with each token represented as an integer index. A note is represented by multiplying the guitar's fret number by the string number to generate a unique index. This approach not only helps the model learn the patterns of movement on the fretboard but also provides flexibility to accommodate alternative tunings, which are increasingly common in

modern guitar playing. With 6 strings and 23 frets per string, the instrument offers 138 distinct fretboard positions. The rhythmic structure is divided into seven beat durations: whole, 1/2, 1/4, 1/8, 1/16, and 1/32 and 1/64 notes. These divisions provide varying levels of temporal granularity for representing note durations. Each beat can further be subdivided into 14 distinct rhythmic values which is listed in Appendix ??, which makes a total to 82 unique beat indices. Since palm muting is a prominent feature in metal music, it is assigned a unique token in the representation. This encoding process produces $138 * 98 * 2 = 27048$ unique musical tokens. The expressive vocabulary listed in Appendix ?? is represented using the same format. A unique integer is assigned to various accents used in guitar playing. The necessity of using relative positional encodings rather than absolute ones, is already shown to enable state reuse without causing temporal confusion ?. Since music is highly sequential such an encoding type should preserve timing, order, and spacing between notes.

Let L be the sequence length and d be the dimensionality of each attention head.

We define the relative positional embedding matrix:

$$\mathbf{R} \in \mathbb{R}^{(2L-1) \times d}$$

where the size $2L-1$ corresponds to all possible relative positions from $-(L-1)$ to $+(L-1)$.

The embedding lookup for a relative position $r = j - i$ (shifted by $L-1$ to be non-negative) is given by:

$$\mathbf{r}(r) = \mathbf{R}[r + (L-1)], \quad r \in \{-(L-1), \dots, 0, \dots, (L-1)\}$$

Here, $\mathbf{r}(r) \in \mathbb{R}^d$ is the embedding vector corresponding to relative position r .

2.1 Note representations

The formula for note encodings is given by:

$$note_encodings = ni + bt * 276 + bb * 276 * 14 \quad (1)$$

$$ni = fn + sn * 23 + 6 * 23 * nval \quad (2)$$

$$nval = \begin{cases} 1 & \text{if, Palm Mute} \\ 0 & \text{otherwise} \end{cases} \quad (3)$$

$$bt = \begin{cases} 1 + d & \text{if, Triplet(3)} \\ 2 + d & \text{if, Quintuplet(5)} \\ 3 + d & \text{if, Sextuplet(6)} \\ 4 + d & \text{if, Septuplet(7)} \\ 5 + d & \text{if, Nonuplet(9)} \\ 6 + d & \text{if, Undecuplet(11)} \end{cases} \quad (4)$$

$$d = \begin{cases} 7 & \text{if, Dotted note} \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

$$bb = \begin{cases} 0 & \text{if, 1 / 64} \\ 1 & \text{if, 1 / 32} \\ 2 & \text{if, 1 / 16} \\ 3 & \text{if, 1 / 8} \\ 4 & \text{if, 1 / 4} \\ 5 & \text{if, 1 / 2} \\ 6 & \text{if, 1 / 1} \end{cases} \quad (6)$$

2.2 Accent representations

2.2.1 CHORD

A chord is a group of notes played together, typically three or more. Figure ?? depicts a C major chord comprising of 3 notes - root note [C] played on the fret 2 / 5th string , major third [G] on fret 1 / 4th string and perfect fifth [E] on fret 2 / 1st string.




Figure 1: C major chord shape on a 6 string with custom tuning.

15413	27051	15389	27051	15321	27051
C	BARRE.NOTE	E	BARRE.NOTE	G	BARRE.NOTE

Table 1: C major chord quarter note tokenisation in a sequence.

2.2.2 OTHER

Accentuations like BEND_NOTE, SLIDE_NOTE etc are represented by the note encodings from Section ?? followed by an accent token listed in Appendix ??.

2.3 EOS and BAR

A music bar (also called a measure) is a fundamental unit in musical notation that divides music into segments of equal time, defined by the time signature. Each bar is separated by a BAR token. A collection of measures making a song is separated by an EOS token.




Figure 2: Fret 5 string 5 quarter note bend

15416	27052
C	BEND_NOTE_1

Table 2: Bend note tokenisation

3 Embedding

A bar consists of a fixed number of beats, composed of musical notes that were tokenised as described in the previous section. Each bar is treated like a sentence. The tokens representing the accents or expression in riff are placed next to each note in the sentence. An example of a riff and its tokenisation is given in Figure ?? . It has been demonstrated that the transformer equipped with relative attention is very well-suited for generative modelling of symbolic music. ?. The chunking strategy used to carry over context between sequences during training follows the approach introduced in ?, which builds upon Transformer-XL’s segment-level recurrence mechanism. This method enables the model to capture longer-range dependencies by reusing hidden states from previous token chunks across segments.

4 Training

4.1 Model

The Transformer is built from a decoder-only stack with masked self attention as described in ?. Although, segment-level recurrence mechanism is captured in Transformer-XL ?, its memory is limited to one previous segment at a time. To model dependencies across a whole batch, the gaussian distributions of the sequences are learned through a Multi layer Perceptron layer. If the input is a batch of sequences defined by:

$$X \in \mathbb{R}^{B \times T \times d} \quad (7)$$

Where:

- B : batch size (number of sequences)
- T : number of tokens per sequence
- d : embedding dimension
- $X_i \in \mathbb{R}^{T \times d}$: the i -th sequence in the batch

Sequence level mean pooling is performed by:

$$\bar{x}_i = \frac{1}{T} \sum_{t=1}^T X_{i,t} \in \mathbb{R}^d \quad (8)$$

Gaussian parameters are then computed with a Multi Layer Perceptron. Let $f : \mathbb{R}^d \rightarrow \mathbb{R}^{2d}$ be a learned MLP. Then:

$$[\mu_i, \log \sigma_i] = f(\bar{x}_i) \quad \text{with} \quad \mu_i, \log \sigma_i \in \mathbb{R}^d \quad (9)$$

Ensure positivity of standard deviation:

$$\sigma_i = \exp(\log \sigma_i) + \varepsilon \quad \text{where} \quad \varepsilon > 0 \quad (10)$$

Each sequence is then represented as a diagonal multivariate Gaussian:

$$\mathcal{N}_i = \mathcal{N}(\mu_i, \text{diag}(\sigma_i^2)) \quad (11)$$

Vectorized for the batch:

$$[\mu, \log \sigma] = f(\bar{X}) \in \mathbb{R}^{B \times 2d} \quad (12)$$

$$\sigma = \exp(\log \sigma) + \varepsilon \in \mathbb{R}^{B \times d} \quad (13)$$

The pairwise KL divergence between each pair of Gaussians ?:

$$\text{KL}(\mathcal{N}_i \parallel \mathcal{N}_j) = \frac{1}{2} \sum_{k=1}^d \left[\log \left(\frac{\sigma_{j,k}^2}{\sigma_{i,k}^2} \right) + \frac{\sigma_{i,k}^2}{\sigma_{j,k}^2} + \frac{(\mu_{i,k} - \mu_{j,k})^2}{\sigma_{j,k}^2} - 1 \right] \quad (14)$$

This produces a KL matrix:

$$\text{KL_matrix} \in \mathbb{R}^{B \times B}, \quad [\text{KL_matrix}]_{i,j} = \text{KL}(\mathcal{N}_i \parallel \mathcal{N}_j) \quad (15)$$

A simpler version of the attention matrix from ? is used from which the KL divergence-based sequence bias is subtracted:

$$\text{RA}_{\text{KL}} = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_k}} + QR^\top - \alpha \cdot \text{KL}_{\text{bias}} \right) \quad (16)$$

A weighted sum over relative positional value contributions V_{rel} is then added:

$$\text{Output} = \text{RA}_{\text{KL}} \times V + \text{RA}_{\text{KL}} \times V_{\text{rel}} \quad (17)$$

4.2 Loss function

The model's loss is composed of three components:

$$\text{Totalloss} = \text{CrossEntropy} + \text{Repetitionloss} + \text{Sequenceloss} \quad (18)$$

4.2.1 CROSS ENTROPY LOSS

Cross-entropy loss is the standard choice for training deep neural networks and logistic regression models for both binary and multi-class classification tasks. ?

4.2.2 REPETITION LOSS

Repetition helps organise musical ideas, making pieces more coherent and memorable. It gives listeners reference points and expectations, aiding in the recognition and recall of themes or motifs. ? Composers have long used stochastic (probability-based) methods to inject randomness into repeated motifs, sometimes even employing quantum randomness for truly unpredictable results. ? This demands careful consideration, as it relies on the listener’s inherent ability to recognise patterns and their natural inclination toward novelty. Since it is already known that tokens listed in Appendix A cannot repeat, it is penalised through repetition loss making them less likely to be selected again in subsequent steps. A variant of the repetition penalty introduced in ? is employed; however, instead of greedy decoding, multinomial sampling is used to generate tokens to get semantically different output sequences ?. Including this loss during training eliminates the need to explicitly handle constraints for generating valid sequences at inference time

$$\text{Let } p_{b,t,v} = \text{softmax}\left(\frac{\text{logits}_{b,t,v}}{T}\right), \quad \text{for } b = 1, \dots, B; \ t = 0, \dots, T-1; \ v = 1, \dots, V \quad (19)$$

$$(20)$$

Compute average past probabilities:

$$\bar{p}_{b,t,v} = \frac{1}{t} \sum_{k=0}^{t-1} p_{b,k,v}, \quad \text{for } t \geq 1 \quad (21)$$

Repetition score (dot product of current and average previous probabilities)

$$R_{b,t} = \sum_{v=1}^V p_{b,t,v} \cdot \bar{p}_{b,t,v}, \quad \text{for } t \geq 1 \quad (22)$$

Final penalty (mean repetition score across batch and sequence, excluding)

$$\text{Repetitionloss} = \frac{\lambda}{B(T-1)} \sum_{b=1}^B \sum_{t=1}^{T-1} R_{b,t} \quad (23)$$

where λ is a repetition penalty weight hyper parameter.

4.2.3 SEQUENCE PENALTY LOSS

Pairwise KL divergence between each pair of Gaussians is given by (??). Given a batch of B sequences, Sequence loss is defined by:

$$\mathcal{L}_{\text{KL}} = \frac{1}{B} \sum_{b=1}^B \text{KL}(\mathcal{N}_q^{(b)} \parallel \mathcal{N}_k^{(b)}) \quad (24)$$

To make the training more stable, the sequence loss is annealed and clamped.

$$\text{Sequenceloss} = \min\left(\frac{\text{step}}{\text{EPOCHS}} \cdot \left(\frac{\delta \cdot \mathcal{L}_{\text{KL}}}{\text{BATCH}}\right), 0.2\right) \quad (25)$$

where, δ is a sequence weight hyper parameter.

5 Inference

Given that the model is autoregressive and expects an initial token, selecting the start token based on the distribution of starting notes of musical measures in the training data yields more consistent generation.




Figure 3: Token occurrences of first note of a musical measure in training data.




Figure 4: Accent occurrences in training data.

The idea of context carryover has been well-established in ?. During inference, the model carries over context from previous sequences by using the last token of the preceding measure as the first token for the subsequent measure.

Loss values observed after 100 epochs of training:

Loss	Value
Cross Entropy	1.4629
Repetition	0.1368
Sequence	0.1985
Total	1.7983
Validation	2.0622

Table 3: Loss values

To compare the quality of predicted tokens (P) from true distributions (Q), relative entropy is calculated using Kullback–Leibler divergence:

$$D_{KL}(P \parallel Q) = \sum_{x \in X} P(x) \log \left(\frac{P(x)}{Q(x)} \right) \quad (26)$$




Figure 5: Note and beat occurrences of training data and inference

The model produces musically pleasing sequences, effectively combining notes, beats, and accents while maintaining sequence context between measures.

Custom Tuning
 ① = F ④ = D#
 ② = C ⑤ = A#
 ③ = G# ⑥ = D#

Moderate $\text{♩} = 120$

Figure 6 shows four musical measures of a guitar riff. Each measure is represented by a staff of music and a corresponding guitar tablature. The first measure is marked with a '1' and 'dist.guit.'. The second measure is marked with a '2' and 'w/bar'. The third measure is marked with a '3' and 'H P'. The fourth measure is marked with a '4' and 'st.'. The tablature includes fret numbers (0-5) and techniques like P.M. (Palm Mute), full, w/bar (with bar), H (Harmonics), P (Palm Mute), and st. (Sustained). The tempo is set to Moderate with a quarter note equal to 120 beats per minute.

Figure 6: Inference results showing 4 musical measures

The model implementation is available on GitHub ?.

Acknowledgments and Disclosure of Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Appendix A.

In this appendix a comprehensive list of accents played on the guitar is provided.

Accent	Unique Index
BAR	27049
BOS	27050
BARRE_NOTE	27051
BEND_NOTE_1	27052
BEND_NOTE_2	27053
BEND_NOTE_3	27054
BEND_NOTE_4	27055
BEND_NOTE_5	27056
BEND_NOTE_6	27057
BEND_NOTE_7	27058
TREM_BAR_1	27059
TREM_BAR_2	27060
TREM_BAR_3	27061
TREM_BAR_4	27062
TREM_BAR_5	27063
DEAD_NOTE	27064
SLIDE_NOTE_1	27065
SLIDE_NOTE_2	27066
SLIDE_NOTE_3	27068
SLIDE_NOTE_4	27069
SLIDE_NOTE_5	27070
SLIDE_NOTE_6	27071
HAMMER	27072
VIBRATO	27073
HARMONIC_1	27074
EOS	27075

Table 4: Example of unique indexing of guitar accents.

Appendix B.

In this appendix a comprehensive list of rhythmic patterns are listed.

Pattern	Unique Index
Base	1
Triplet	2
Quintuplet	3
sextuplet	4
septuplet	5
9 tuplets	6
11-tuplets	7
Dotted - Base	8
Dotted - Triplet	9
Dotted - Quintuplet	10
Dotted - sextuplet	11
Dotted - septuplet	12
Dotted - 9 tuplets	13
Dotted - 11-tuplets	14

Table 5: Example of unique indexing of rhythm.